

The logo for Devoxx is located on the left side of the slide. It consists of the word "Devoxx" in a stylized, white, blocky font on a black background. The "x"s are orange with white outlines. A small "TM" trademark symbol is at the top left of the logo.

# Nuts and Bolts of WebSocket


Arun Gupta, Red Hat, @arungupta  
Peter Moskovits, Kaazing, @peterm\_kaazing



# Agenda

- Introduction
- WebSocket and Node.js
- WebSocket using JSR 356
  - Server
  - Client
- Securing WebSocket
- Embedded WebSocket
- Load Balance WebSocket
- Pub/Sub over WebSocket
  - STOMP over WebSocket
  - MQTT over WebSocket
- REST and SSE
- Scalability
- Debugging
- Production Tips

# The “long” story of WebSocket





# “Limitations” of HTTP

- Client-driven
- Half-duplex
- Verbose
- New TCP connection



# “Hello World” HTTP

```
POST /websocket-vs-rest-payload/webresources/rest HTTP/1.1\r\nHost: localhost:8080\r\nConnection: keep-alive\r\nContent-Length: 11\r\nUser-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_1) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/32.0.1700.107 Safari/537.36\r\nOrigin: chrome-extension://hgmlloofddffdnphfgcellkdfbfjeloo\r\nContent-Type: text/plain \r\nAccept: */*\r\nAccept-Encoding: gzip,deflate,sdch\r\nAccept-Language: en-US,en;q=0.8\r\n\r\n
```

```
HTTP/1.1 200 OK\r\nConnection: keep-alive\r\nX-Powered-By: Undertow 1\r\nServer: Wildfly 8 \r\nContent-Type: text/plain\r\nContent-Length: 11 \r\nDate: Fri, 21 Feb 2014 21:27:53 GMT \r\n\r\n
```

663 bytes



# How WebSocket solves it ?

- Bi-directional (client-driven)
- Full-duplex (half-duplex)
- Lean protocol (verbose)
- Single TCP connection (new TCP)



# What is WebSocket ?

- Bi-directional, full-duplex, communication channel over a single TCP connection
- Originally proposed as part of HTML5
- IETF-defined **Protocol**: RFC 6455
- W3C-defined **JavaScript API**

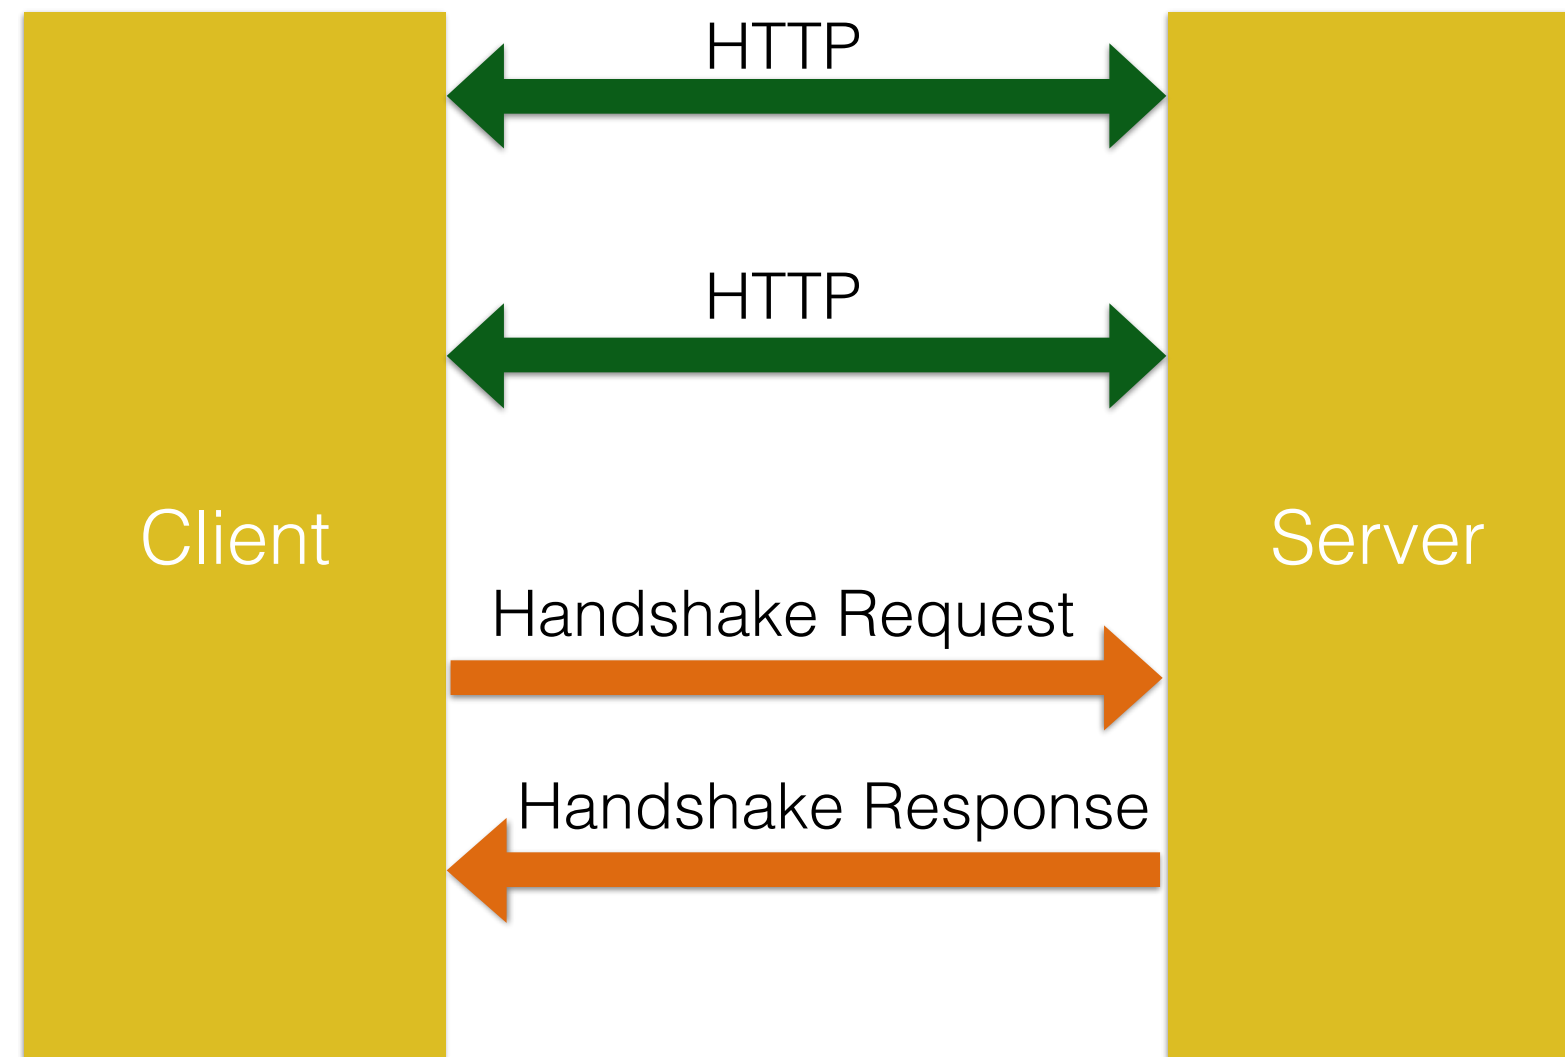


# How does it work ?

- Upgrade HTTP to WebSocket (single TCP connection)
- Send data frames in both direction (bi-directional)
- Send messages independent of each other (full-duplex)
- End the connection



# How does it work ?





# Handshake Request

GET /chat HTTP/1.1

Host: server.example.com

Upgrade: websocket

Connection: Upgrade

Origin: http://example.com

Sec-WebSocket-Key: dGhlIHNhbXBsZSBub25jZQ==

Sec-WebSocket-Protocol: chat, superchat

Sec-WebSocket-Version: 13



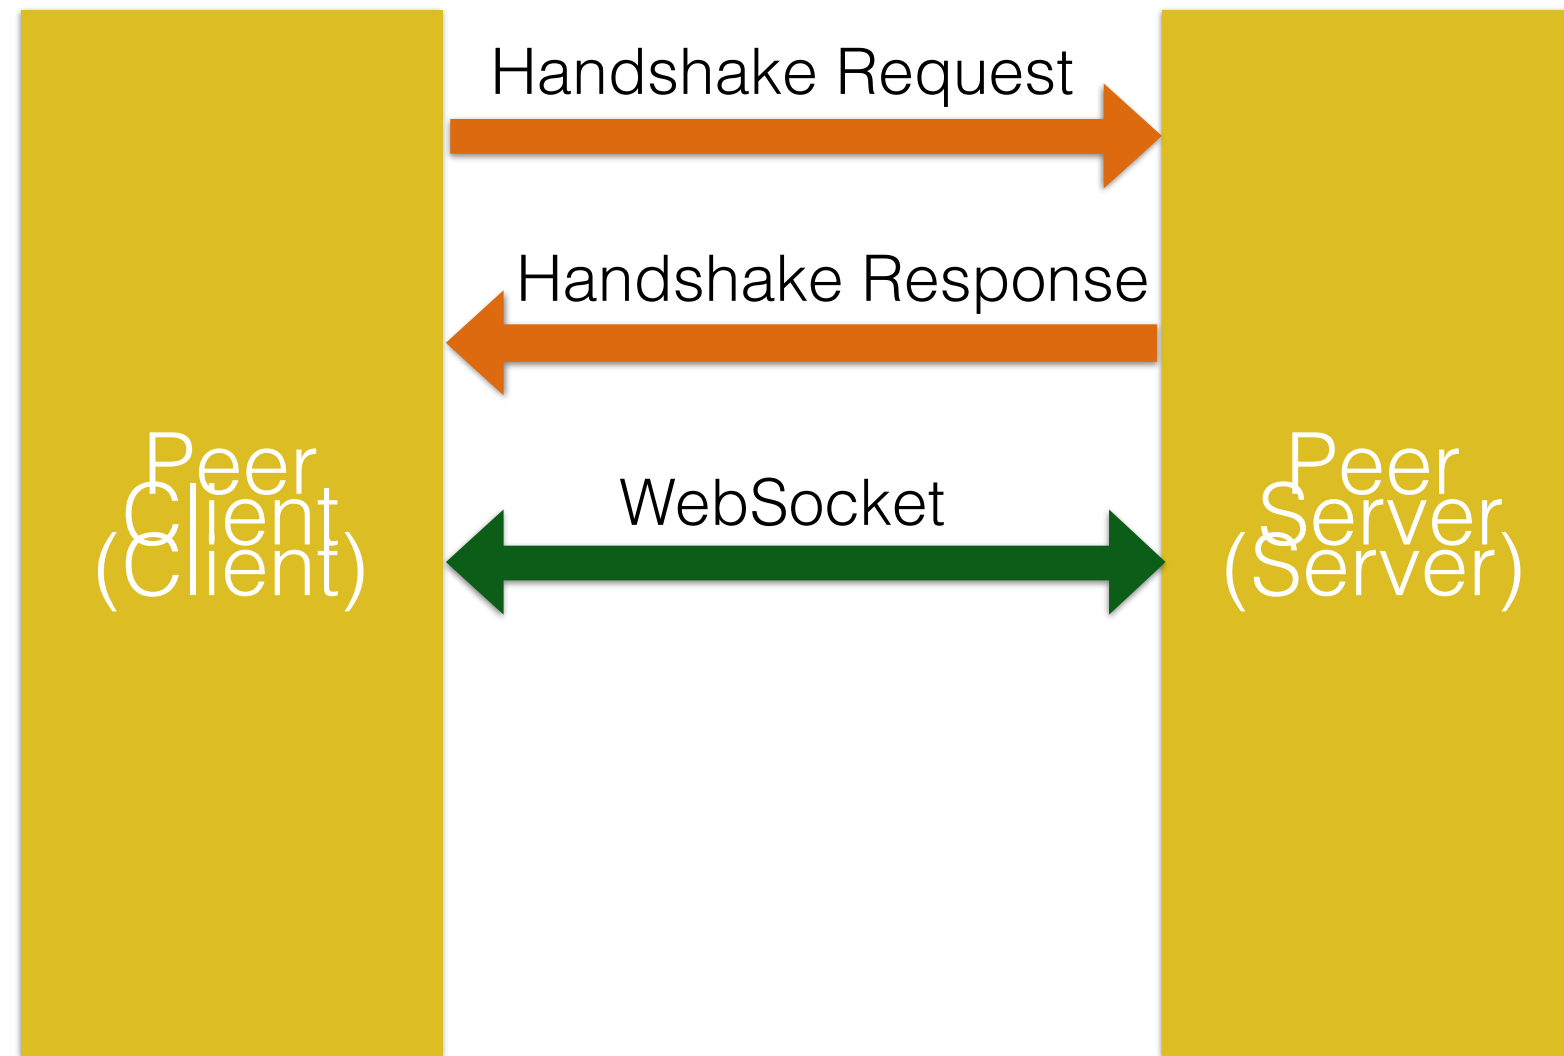
# Handshake Response

HTTP/1.1 101 Switching Protocols

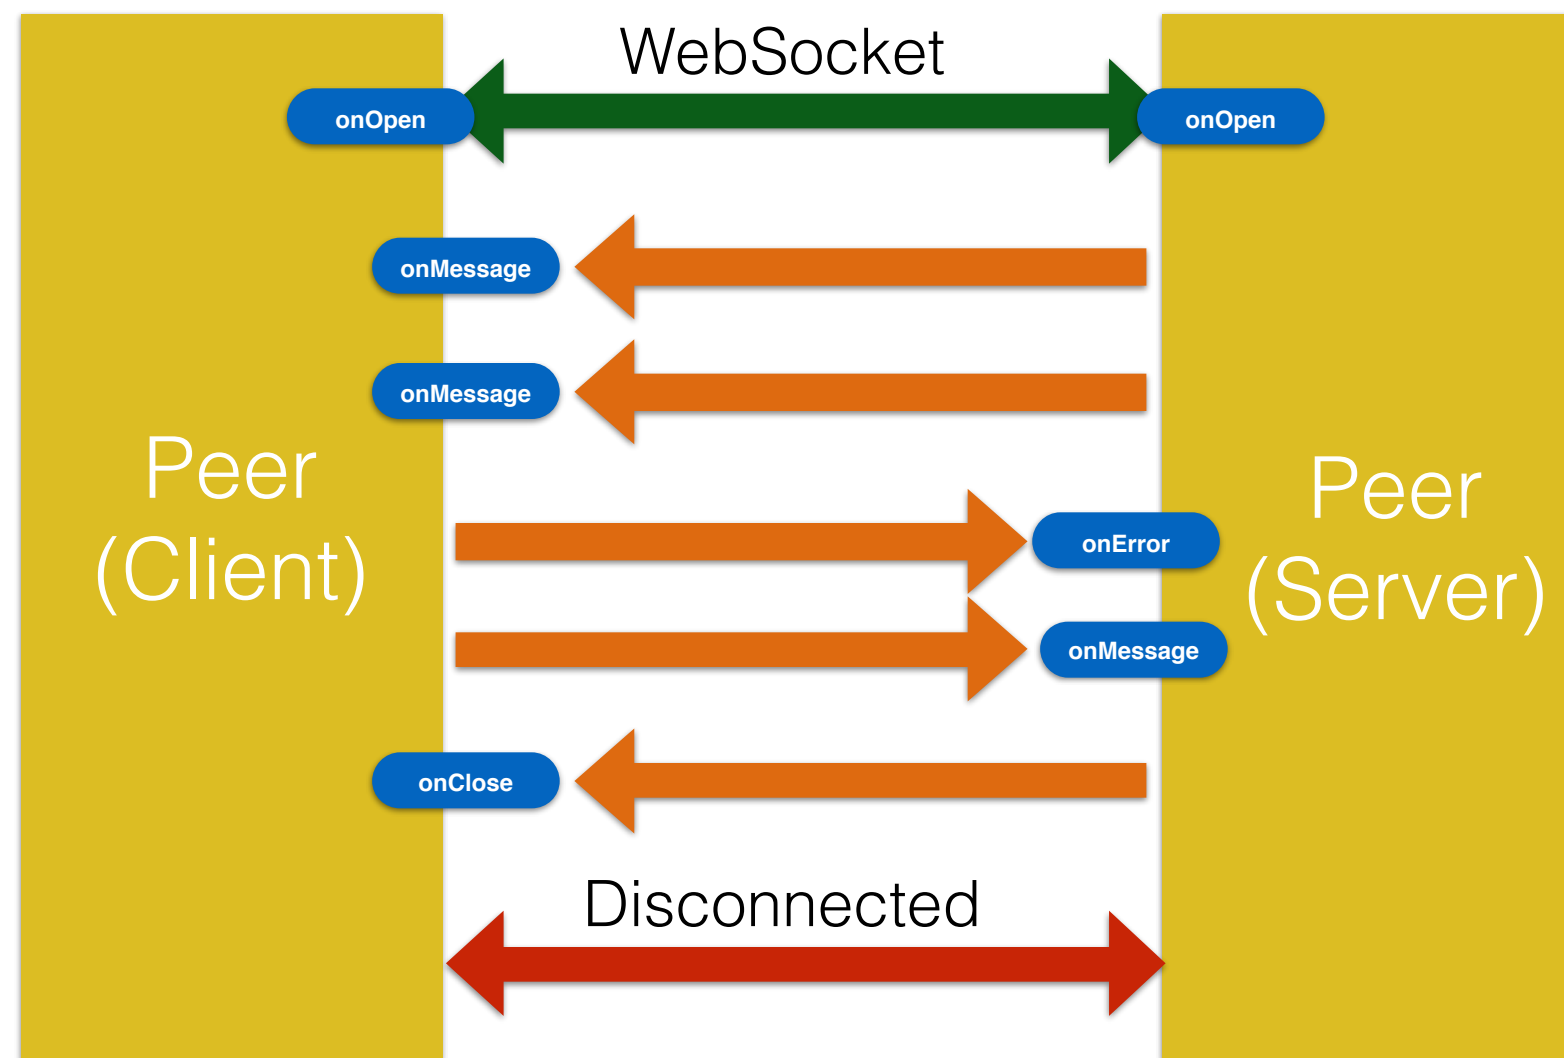
Upgrade: websocket  
Connection: Upgrade

Sec-WebSocket-Accept: s3pPLMBiTxaQ9kYGzzhZRbK+xOo=  
Sec-WebSocket-Protocol: chat

# How does it work ?



# How does it work ?





# WebSocket Subprotocols

- Facilitates application layer protocols
- Registered in a Subprotocol name registry
  - Identifier, common name, definition
  - [www.iana.org/assignments/websocket/websocket.xml#subprotocol-name](http://www.iana.org/assignments/websocket/websocket.xml#subprotocol-name)
  - STOMP, XMPP, MQTT, SOAP, ...



# WebSocket Extensions

- Add capabilities to the base protocol
- Multiplexing  
<http://tools.ietf.org/html/draft-tamplin-hybi-google-mux>
- Compression: Only non-control frames/messages
  - Per-frame  
<http://tools.ietf.org/html/draft-tyoshino-hybi-websocket-perframe-deflate>
  - Per-message  
<http://tools.ietf.org/html/draft-ietf-hybi-permessage-compression>



# WebSocket JavaScript API

```
[Constructor(DOMString url, optional (DOMString or DOMString[]) protocols)]
interface WebSocket : EventTarget {
  readonly attribute DOMString url;

  // ready state
  const unsigned short CONNECTING = 0;
  const unsigned short OPEN = 1;
  const unsigned short CLOSING = 2;
  const unsigned short CLOSED = 3;
  readonly attribute unsigned short readyState;
  readonly attribute unsigned long bufferedAmount;

  // networking
  attribute EventHandler onopen;
  attribute EventHandler onerror;
  attribute EventHandler onclose;
  readonly attribute DOMString extensions;
  readonly attribute DOMString protocol;
  void close([Clamp] optional unsigned short code, optional DOMString reason);

  // messaging
  attribute EventHandler onmessage;
  attribute DOMString binaryType;
  void send(DOMString data);
  void send(Blob data);
  void send(ArrayBuffer data);
  void send(ArrayBufferView data);
};
```

[www.w3.org/TR/websockets](http://www.w3.org/TR/websockets)







# Developer Tools

DEVONX<sup>TM</sup>





# Java API for WebSocket

- API for WebSocket server and client endpoint
  - Annotated: `@ServerEndpoint`, `@ClientEndpoint`
  - Programmatic: `Endpoint`
    - WebSocket opening handshake negotiation
- Lifecycle Callback methods
- Integration with Java EE technologies



# Annotated Endpoint

```
import javax.websocket.*;  
  
@ServerEndpoint("/hello")  
public class HelloBean {  
    @OnMessage  
    public String sayHello(String name) {  
        return "Hello " + name;  
    }  
}
```



# WebSocket Annotations

- Class-level annotations
  - `@ServerEndpoint`: Turns a POJO in a server endpoint
  - `@ClientEndpoint`: Turns a POJO in a client endpoint



# WebSocket Annotations

- Method-level annotations
  - `@OnMessage`: Intercepts WebSocket messages
  - `@OnOpen`: Intercepts WebSocket open events
  - `@OnClose`: Intercepts WebSocket close events
  - `@OnError`: Intercepts WebSocket error events



# WebSocket Annotations

- Parameter-level annotations
  - `@PathParam`: Matches path segment of a URI-template





# @ServerEndpoint attributes

- `value`: Relative URI or URI template e.g. `‘/hello’` or `‘/chat/{subscriber-level}’`
- `decoders`: list of message decoder classnames
- `encoders`: list of message encoder classnames
- `subprotocols`: list of the names of the supported subprotocols



# Chat Server

```
@ServerEndpoint("/chat")
public class ChatBean {
    static Set<Session> peers = Collections.synchronizedSet("...");

    @OnOpen
    public void onOpen(Session peer) {
        peers.add(peer);
    }

    @OnClose
    public void onClose(Session peer) {
        peers.remove(peer);
    }

    @OnMessage
    public void message(String message) {
        for (Session peer : peers) {
            peer.getBasicRemote().sendObject(message);
        }
    }
}
```



# Chat Server Simplified

```
@ServerEndpoint("/chat")
public class ChatBean {
    @OnMessage
    public void message(String message, Session endpoint) {
        for (Session peer : endpoint.getOpenSessions()) {
            peer.getBasicRemote().sendObject(message);
        }
    }
}
```



Hello there!

Howdy?

<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/chat>



# Custom Payloads

```
@ServerEndpoint(  
    value="/hello",  
    decoders={MyMessageDecoder.class},  
    encoders={MyMessageEncoder.class}  
)  
public class MyEndpoint {  
    . . .  
}
```



# Custom Payloads: Text decoder

```
public class MyMessageDecoder implements Decoder.Text<MyMessage> {  
    public MyMessage decode(String s) {  
        JsonObject jsonObject = Json.createReader("...").readObject();  
        return new MyMessage(jsonObject);  
    }  
  
    public boolean willDecode(String string) {  
        . . .  
        return true;  
    }  
  
    . . .  
}
```



# Custom Payloads: Text encoder

```
public class MyMessageDecoder implements Encoder.Text<MyMessage> {  
    public String encode(MyMessage myMessage) {  
        return myMessage.jsonObject.toString();  
    }  
    . . .  
}
```



# Custom Payloads: Binary decoder

```
public class MyMessageDecoder implements Decoder.Binary<MyMessage> {  
    public MyMessage decode(byte[] s) {  
        . . .  
        return myMessage;  
    }  
  
    public boolean willDecode(byte[] string) {  
        . . .  
        return true;  
    }  
  
    . . .  
}
```







# Method Signatures

- Exactly one of the following
  - **Text:** `String`, `boolean`, Java primitive or equivalent class, `Reader`, any type for which there is a decoder
  - **Binary:** `byte[ ]`, `ByteBuffer`, `byte[ ]` and `boolean`, `ByteBuffer` and `boolean`, `InputStream`, any type for which there is a decoder
  - **Pong messages:** `PongMessage`
- An optional **Session** parameter
- 0..n String parameters annotated with **@PathParam**



# Sample Messages

- `void m(String s);`
- `void m(Float f, @PathParam("id")int id);`
- `Product m(Reader reader, Session s);`
- `void m(byte[] b);` or `void m(ByteBuffer b);`
- `Book m(int i, Session s, @PathParam("isbn")String isbn, @PathParam("store")String store);`



<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/whiteboard>



# URI Template Matching

```
@ServerEndpoint("/chat/{roomId}")
public class ChatServer {
    @OnMessage
    public void receiveMessage(
        @PathParam("roomId")String roomId) {
        . . .
    }
}
```





# Client Endpoint

**@ClientEndpoint**

```
public class HelloClient {  
    @OnMessage public void message(  
        String message,  
        Session session) {  
        // . . .  
    }  
}
```

```
WebSocketContainer c = ContainerProvider.getWebSocketContainer();  
c.connectToServer(HelloClient.class, "hello");
```



<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/google-docs>



# Programmatic Endpoint

```
public class ChatServer extends Endpoint {  
    @Override  
    public void onOpen(Session session) {  
        session.addMessageHandler(new MessageHandler.Text() {  
            public void onMessage(String message) {  
                try {  
                    session  
                        .getBasicRemote()  
                        .sendText(message);  
                } catch (IOException ex) { }  
            }  
        });  
    }  
}
```



# Programmatic Endpoint Config

```
public class MyEndpointConfig implements ServerApplicationConfig {

    @Override
    public Set<ServerEndpointConfig> getEndpointConfigs(
        Set<Class<? extends Endpoint>> set) {
        return new HashSet<ServerEndpointConfig>() {
            {
                add(ServerEndpointConfig.Builder
                    .create(ChatServer.class, "/chat")
                    .build());
            }
        };
    }

    @Override
    public Set<Class<?>> getAnnotatedEndpointClasses(Set<Class<?>> set) {
        return Collections.emptySet();
    }
}
```





<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/endpoint-singleton>



# Securing WebSockets

- Origin-based security model
- Sec-xxx keys can not be set using XMLHttpRequest
  - Sec-WebSocket-Key, Sec-WebSocket-Version
- User-based security using Servlet security mechanism
  - Endpoint mapped by **ws://** is protected using security model defined using the corresponding http:// URI
  - Authorization defined using `<security-constraint>`
- Transport Confidentiality using **wss://**
  - Access allowed over encrypted connection only



# Cross-Origin Resource Sharing

- Relaxes same-origin restrictions to network requests
- Servers include `Access-Control-Allow-Origin` HTTP header
- `Access-Control-*` headers
  - `Max-Age`
  - `Allow-Methods`
  - `Allow-Headers`
  - ...
- [www.w3.org/TR/cors/](http://www.w3.org/TR/cors/)



# User-based Security

<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/endpoint-security>



# TLS-based Security

<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/endpoint-wss>



# Embedded WebSocket

- Undertow New web server in WildFly 8
- Blocking and non-blocking based on NIO
- Composition/handler-based architecture
- Lightweight and fully embeddable
- Supports Servlet 3.1 and HTTP Upgrade
- mod\_cluster supported



# Undertow is awesome!

```
techempower@lg01:~$ wrk -d 30 -c 256 -t 40 http://10.0.3.2:8080/byte
Running 30s test @ http://10.0.3.2:8080/byte
 40 threads and 256 connections
  Thread Stats   Avg      Stdev     Max   +/-  Stdev
    Latency    247.05us    3.52ms  624.37ms   99.90%
   Req/Sec    27.89k     6.24k   50.22k   71.15%
 31173283 requests in 29.99s  3.83GB read
Socket errors: connect 0, read 0, write 0, timeout 9
Requests/sec: 1039305.27
Transfer/sec:   130.83MB
```

This is output from [Wrk](#) testing a single server running [Undertow](#) using conditions similar to Google's test (1-byte response body, no HTTP pipelining, no special request headers) **1.039 million requests per second.**

<http://www.techempower.com/blog/2014/03/04/one-million-http-rps-without-load-balancing-is-easy/>



`git@github.com:undertow-io/undertow.git`





# Load Balance WebSocket

- Reverse proxy
  - Apache module: `mod_proxy_wstunnel`
- Only vertical scaling
- No session replication



<http://blog.arungupta.me/2014/08/load-balance-websockets-apache-httpd-techtip48/>



# Pub/Sub over WebSocket

<https://github.com/arun-gupta/kaazing-openshift-cartridge>



# STOMP over WebSocket

- STOMP: Simple Text Oriented Messaging Protocol
- Interoperable wire format: any client, any broker
- Messaging interoperability among languages and platforms
  - Unlike JMS
- REST for messaging: CONNECT, SEND, SUBSCRIBE, ...
- Map STOMP frames to WebSocket frames





<https://github.com/arun-gupta/wildfly-samples/tree/master/websocket-stomp>



# MQTT over WebSocket

- Light-weight pub/sub messaging over TCP
- Designed for “small foot print” or limited bandwidth
- MQTT 3.1.1 just became an OASIS standard
- Plain byte array message payload
- Quality-of-Service: 0 (TCP), 1 (at least once), 2 (missed messages)
- “Last will and testament” publish a message after client goes offline





<http://blog.kaazing.com/2013/10/01/controlling-physical-devices-on-the-real-time-web-kaazing-iot-talk-at-javaone-2013/>



# Compare with REST

<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/websocket-vs-rest-payload>

<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/websocket-vs-rest>





# Server-Sent Events

- Part of HTML5 Specification
- Server-push notifications
- Cross-browser JavaScript API: `EventSource`
- Message callbacks
- MIME type: `text/eventstream`

# EventSource API

```
[Constructor(DOMString url, optional EventSourceInit eventSourceInitDict)]
interface EventSource : EventTarget {
  readonly attribute DOMString url;
  readonly attribute boolean withCredentials;

  // ready state
  const unsigned short CONNECTING = 0;
  const unsigned short OPEN = 1;
  const unsigned short CLOSED = 2;
  readonly attribute unsigned short readyState;

  // networking
  [TreatNonCallableAsNull] attribute Function? onopen;
  [TreatNonCallableAsNull] attribute Function? onmessage;
  [TreatNonCallableAsNull] attribute Function? onerror;
  void close();
};

dictionary EventSourceInit {
  boolean withCredentials = false;
};
```



# WebSockets and SSE ?

| WebSocket                       | Server-Sent Event                   |
|---------------------------------|-------------------------------------|
| Over a custom protocol          | Over simple HTTP                    |
| Full-duplex, bi-directional     | Server-push only, client-server OOB |
| Native support in most browsers | Can be poly-filled to backport      |
| Not straight forward protocol   | Simpler protocol                    |



# WebSockets and SSE ?

| WebSocket                                  | Server-Sent Event                              |
|--|--|
| Application-specific reconnection          | Built-in support for reconnection and event id |
| Require server and/or proxy configurations | No server or proxy change required             |
| Text and Binary                            | Text only                                      |
| Pre-defined message handlers               | Pre-defined and arbitrary                      |




# What makes them scalable ?

- No HTTP/TCP opening/closing connections
  - Handshake over a single TCP connection
  - HTTP connections have short connection timeout (5 secs for Apache)
- Elimination of HTTP headers (cookies, content-type, user-agent, ...)
  - Reduces bandwidth dramatically



# What makes them scalable ?

- Minimal data framing
  - 2-14 bytes overhead after handshake
- Maintaining a TCP connection on server is relatively inexpensive
- Smaller data fragments can be sent without out a request/response
  - Live pushes, e.g. stock sticker
  - Lower latency







# What makes them scalable ?

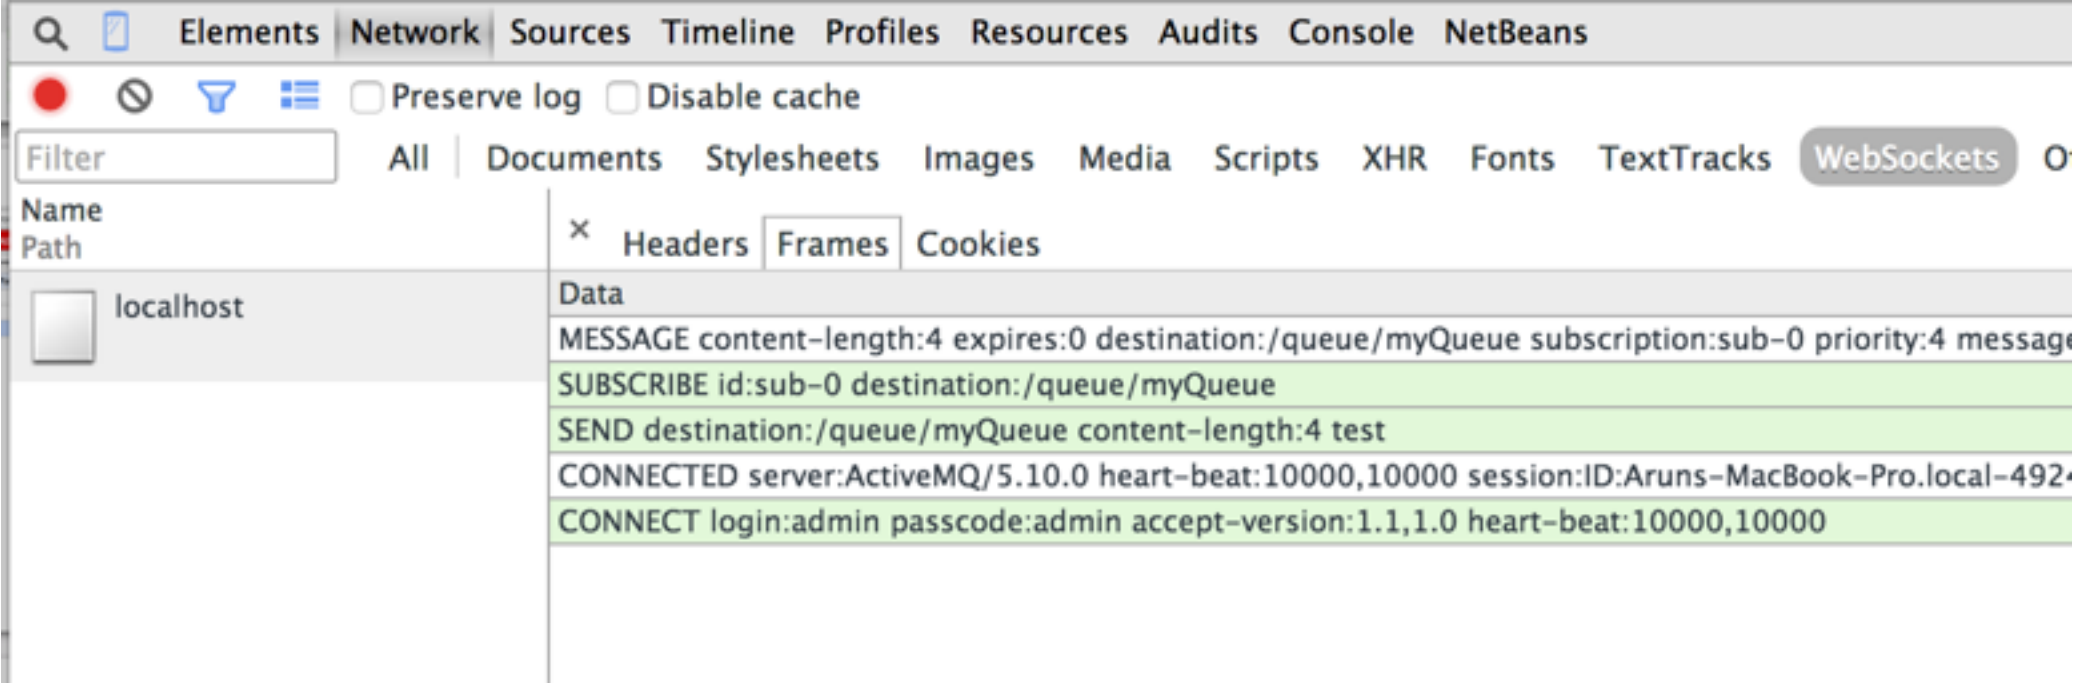
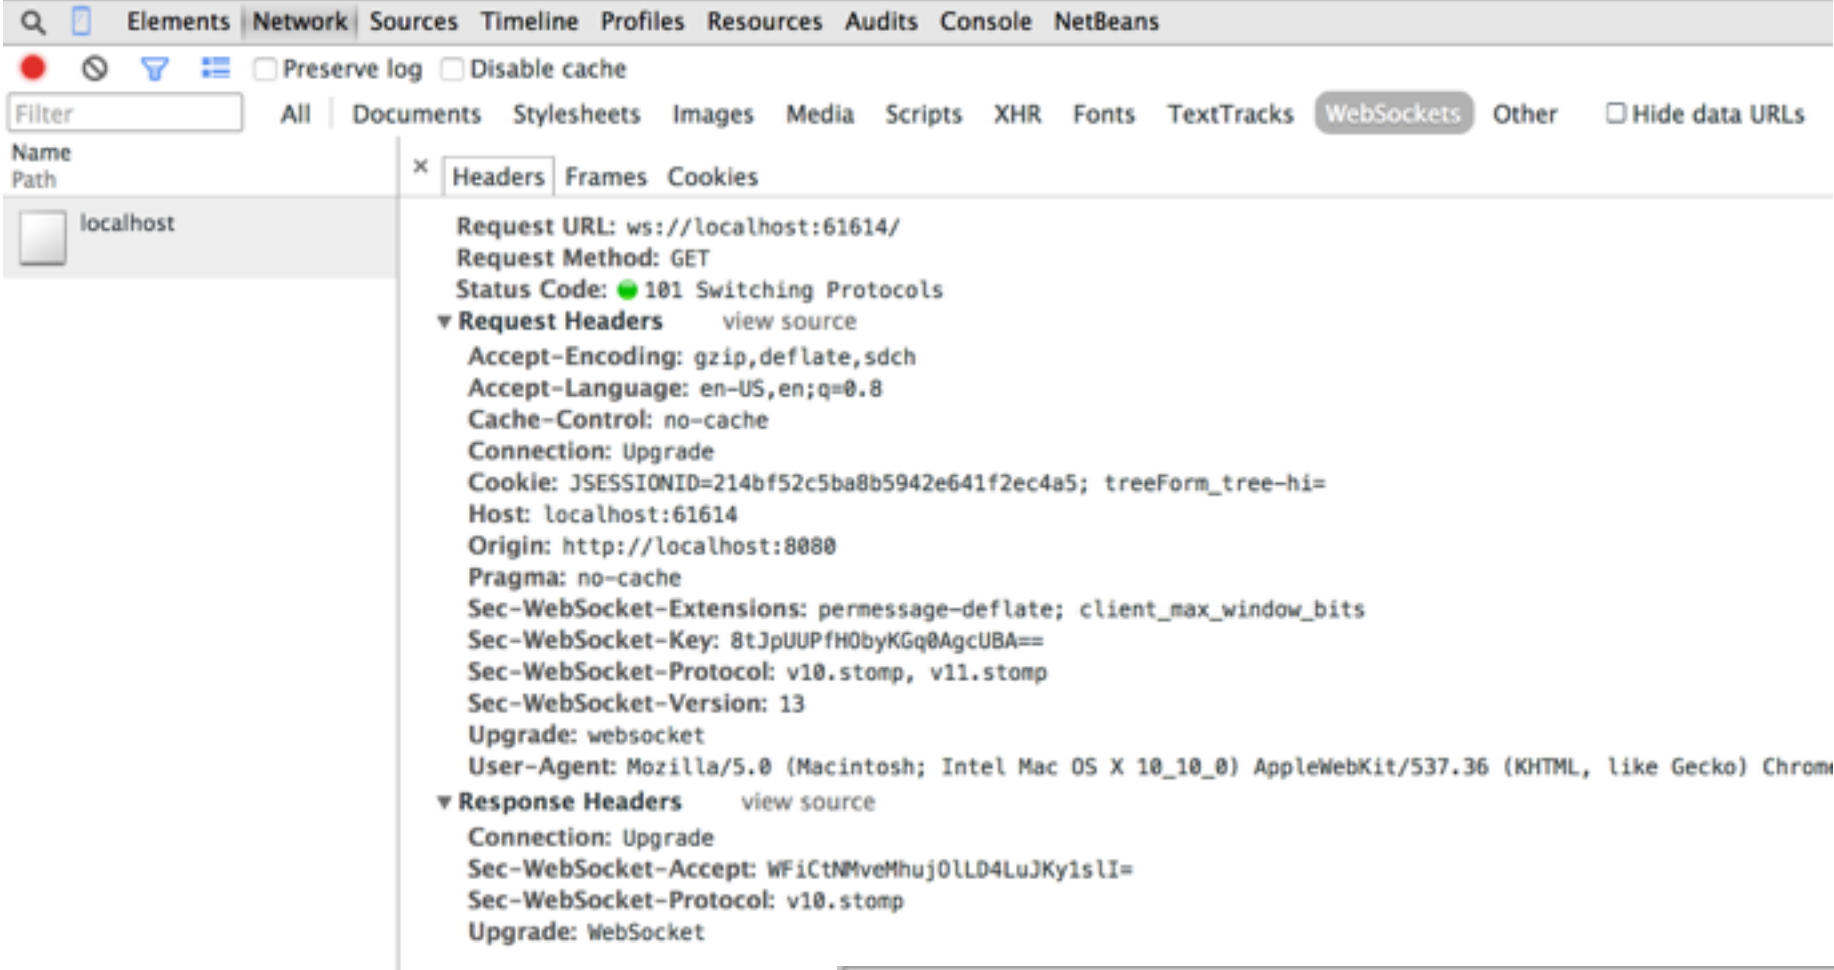
- WebSockets are good at scaling vertically
- HTTP servers are typically configured to log start/completion of HTTP request, not so for WebSocket
- Polling and Long-polling is a waste of bandwidth, WebSockets are more elegant
- Number of concurrent clients depend upon FD settings

TM



# Debugging WebSockets







chrome://net-internals/#events&q=type:SOCKET%20is:active

Events

capturing events (4641)

(?) type:SOCKET is:active

3 of 572

| <input type="checkbox"/>            | ID      | Source Type | Description     |
|-------------------------------------|---------|-------------|-----------------|
| <input type="checkbox"/>            | 1596087 | SOCKET      | localhost:8080  |
| <input type="checkbox"/>            | 1596088 | SOCKET      | localhost:8080  |
| <input checked="" type="checkbox"/> | 1596381 | SOCKET      | localhost:61614 |

1596381: SOCKET

localhost:61614

Start Time: 2014-10-23 10:30:37.952

```

t=143883 [st= 0] +SOCKET_ALIVE [dt=?]
--> source_dependency = 1596378 (CONNECT_JOB)
t=143883 [st= 0] +TCP_CONNECT [dt=0]
--> address_list = ["[::1]:61614"]
t=143883 [st= 0] TCP_CONNECT_ATTEMPT [dt=0]
--> address = "[::1]:61614"
t=143883 [st= 0] -TCP_CONNECT
--> source_address = "[::1]:57613"
t=143883 [st= 0] +SOCKET_IN_USE [dt=?]
--> source_dependency = 1596376 (HTTP_STREAM_JOB)
t=143884 [st= 1] SOCKET_BYTES_SENT
--> byte_count = 616
t=143885 [st= 2] SOCKET_BYTES_RECEIVED
--> byte_count = 164
t=143886 [st= 3] SOCKET_BYTES_SENT
--> byte_count = 89
t=143888 [st= 5] SOCKET_BYTES_RECEIVED
--> byte_count = 134
t=144739 [st= 856] SOCKET_BYTES_SENT
--> byte_count = 61
t=147073 [st= 3190] SOCKET_BYTES_SENT
--> byte_count = 54
t=147077 [st= 3194] SOCKET_BYTES_RECEIVED
--> byte_count = 196
t=153901 [st=10018] SOCKET_BYTES_RECEIVED
--> byte_count = 3
t=154226 [st=10343] SOCKET_BYTES_SENT
--> byte_count = 7

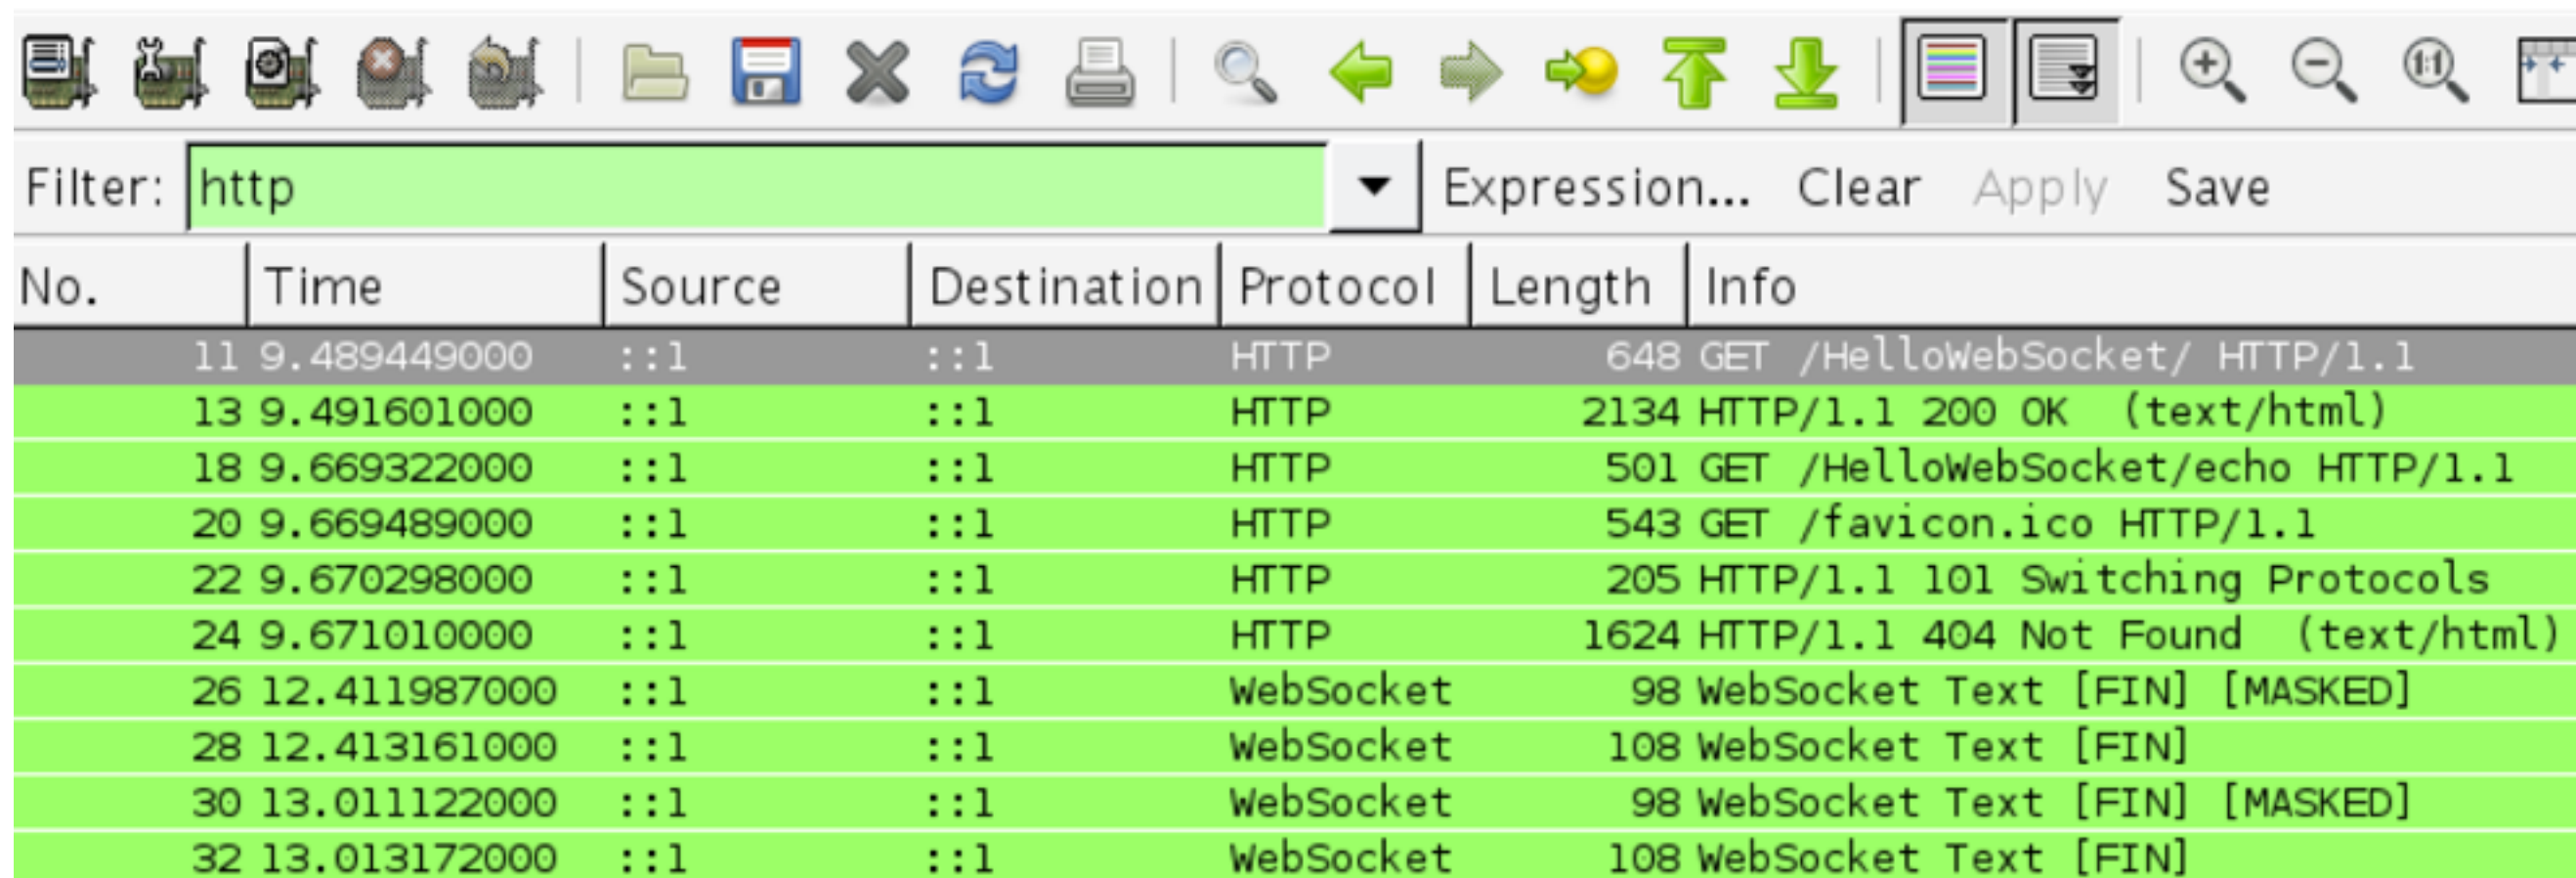
```

#Devoxx #WebSocket

66

@arungupta @peterm\_kaazing

# Debugging WebSockets



| No. | Time         | Source | Destination | Protocol  | Length | Info                               |
|-----|--------------|--------|-------------|-----------|--------|------------------------------------|
| 11  | 9.489449000  | ::1    | ::1         | HTTP      | 648    | GET /HelloWebSocket/ HTTP/1.1      |
| 13  | 9.491601000  | ::1    | ::1         | HTTP      | 2134   | HTTP/1.1 200 OK (text/html)        |
| 18  | 9.669322000  | ::1    | ::1         | HTTP      | 501    | GET /HelloWebSocket/echo HTTP/1.1  |
| 20  | 9.669489000  | ::1    | ::1         | HTTP      | 543    | GET /favicon.ico HTTP/1.1          |
| 22  | 9.670298000  | ::1    | ::1         | HTTP      | 205    | HTTP/1.1 101 Switching Protocols   |
| 24  | 9.671010000  | ::1    | ::1         | HTTP      | 1624   | HTTP/1.1 404 Not Found (text/html) |
| 26  | 12.411987000 | ::1    | ::1         | WebSocket | 98     | WebSocket Text [FIN] [MASKED]      |
| 28  | 12.413161000 | ::1    | ::1         | WebSocket | 108    | WebSocket Text [FIN]               |
| 30  | 13.011122000 | ::1    | ::1         | WebSocket | 98     | WebSocket Text [FIN] [MASKED]      |
| 32  | 13.013172000 | ::1    | ::1         | WebSocket | 108    | WebSocket Text [FIN]               |





# Production Tips

- Proxy can be evil and make WebSockets unusable
  - Issue: Remove “Upgrade” header
  - Fix: Set timeout, remove after onOpen called
  - Issue: Close connection after X idle time
  - Fix: Application-level heartbeat
  - Issue: Not allow to pass through at all
  - Fix: Fall back on long-polling



# Production Tips

- Load Balancer
  - Issue: Don't work with WebSocket, e.g. Amazon ELB
  - Fix: ELB configured to use TCP, but no session affinity
- Browsers
  - Issue: IE 6, 7, 8, 9 and Safari 5 do not support WebSocket
  - Fix: Fallback using Atmosphere, Socket.IO, SockJS
- Inconsistencies in JSR 356



# Atmosphere



- Java/JavaScript framework
- Portable asynchronous applications
- Fallback to long-polling in absence of WebSocket
- Containers: Netty, Jetty, GlassFish, Tomcat, JBoss, WildFly, WebLogic, Resin, WebSphere
- Browsers: Firefox, Chrome, IE (6x+), Opera, Safari, Android

<https://github.com/javaee-samples/javaee7-samples/tree/master/websocket/atmosphere-chat>



# Resources

- Material: [github.com/arun-gupta/nuts-and-bolts-of-websocket](https://github.com/arun-gupta/nuts-and-bolts-of-websocket)